

### Question 10.3

1. Using the GermanCredit data set `germancredit.txt` from <http://archive.ics.uci.edu/ml/machine-learning-databases/statlog/german/> (description at <http://archive.ics.uci.edu/ml/datasets/Statlog+%28German+Credit+Data%29>), use logistic regression to find a good predictive model for whether credit applicants are good credit risks or not. Show your model (factors used and their coefficients), the software output, and the quality of fit. You can use the `glm` function in R. To get a logistic regression (logit) model on data where the response is either zero or one, use `family=binomial(link="logit")` in your `glm` function call.

### Summary :

- German Credit Data set has various data points including binary, numerical and categorical data. The data cleaning was to use `ifelse` method on all the data including the response to binary data points for ease of analysis downstream
- Perform a total of 2 iterations to get to the most significant predictors; eliminating insignificant variables
- Build Logistic Model
- Because of FP cost 5x more than FN, need to fine tune the threshold and the optimal threshold for classification was at 13%
- At last, the AUC at 69% dictates that a random person from the yes group would probably get a higher estimate relative to the person from the no group

## Observation

The test set sensitivity for positive class '1' is higher in the test set than from cross validation on the build set since 0.88 is greater than 0.4018385

Also test set specificity for negative class '0' is lower than from cross validation since 0.4018 is much lower than 0.7988784

This is a good result still since there is a greater cost for misclassifying the positive class as the negative class.

Test set area under the curve is less than Builddata cross validated AUC of 0.69.

```
# ----- Code for Question 10.1.1 -----
```

```
# Clear environment
```

```
rm(list = ls())
```

```
# Setting the random number generator seed so that our results are reproducible
```

```
# (Your solution doesn't need this, but it's usually good practice to do)
```

```
set.seed(1)
```

```
#----- Load Libraries ----- #
```

```
# ----- Data manipulation -----
```

```
#First, Read in the data
```

```
#
```

```
data_gc = read.table("germancredit.txt", sep = "")
```

```
#
```

```
# optional check to make sure the data is read correctly
```

```
#
```

```
head(data_gc,2)
```

```
  V1 V2  V3  V4  V5  V6  V7 V8  V9  V10 V11  V12 V13  V14  V15 V16  V17 V18  V19  V20 V21
1 A11  6 A34 A43 1169 A65 A75  4 A93 A101  4 A121 67 A143 A152  2 A173  1 A192 A201  1
2 A12 48 A32 A43 5951 A61 A73  2 A92 A101  2 A121 22 A143 A152  1 A173  1 A191 A201  2
```

```
#Relabel the response variable as 0 and 1. Set 1 to 0 for 'Good' and 2 to 1 for 'Bad'
```

```
#bad is the positive class
```

```
data_gc$V21<-ifelse(data_gc$V21==1,0,ifelse(data_gc$V21==2,1,data_gc$V21))
```

```
summary(data_gc$V21)
```

```
  Min. 1st Qu.  Median    Mean 3rd Qu.   Max.
```

```
  0.0   0.0   0.0   0.3   1.0   1.0
```

```
prop.table(table(data_gc$V21))
```

```
 0  1
```

```
0.7 0.3
```

```
#Based on above probability 70% Good Credit risk and 30% bad
```

```
#split the data into 70% building,30% test
```

```
#Randomize the data first
```

```
data_gc1 <- data_gc[sample(1:nrow(data_gc)),]
```

```
random_row<- sample(1:nrow(data_gc1),as.integer(0.7*nrow(data_gc1 )),replace=F)
```

```
builddata = data_gc1[random_row,]
```

```
#Assign the test data set to the remaining 30% of the original set
```

```
testdata = data_gc1 [-random_row,]
```

```
table(data_gc1$V21)
```

#Do glm on all future set in Build Data

```
set.seed(713)
```

```
reg1 <- glm(V21 ~.,family=binomial(link = "logit"),data=bulddata)
```

```
summary(reg1 )
```

**Call:**

```
glm(formula = V21 ~ ., family = binomial(link = "logit"), data = bulddata)
```

**Deviance Residuals:**

| Min     | 1Q      | Median  | 3Q     | Max    |
|---------|---------|---------|--------|--------|
| -2.0550 | -0.6914 | -0.3713 | 0.6255 | 2.7030 |

**Coefficients:**

|             | Estimate   | Std. Error | z value | Pr(> z )     |
|-------------|------------|------------|---------|--------------|
| (Intercept) | 1.281e+00  | 1.405e+00  | 0.911   | 0.362054     |
| V1A12       | -2.793e-01 | 2.673e-01  | -1.045  | 0.296080     |
| V1A13       | -1.580e+00 | 5.365e-01  | -2.945  | 0.003225 **  |
| V1A14       | -1.563e+00 | 2.790e-01  | -5.601  | 2.14e-08 *** |
| V2          | 3.432e-02  | 1.139e-02  | 3.014   | 0.002575 **  |
| V3A31       | 3.189e-02  | 6.695e-01  | 0.048   | 0.962006     |
| V3A32       | -6.360e-01 | 5.097e-01  | -1.248  | 0.212157     |
| V3A33       | -5.165e-01 | 5.570e-01  | -0.927  | 0.353753     |
| V3A34       | -1.239e+00 | 5.128e-01  | -2.417  | 0.015640 *   |
| V4A41       | -1.500e+00 | 4.529e-01  | -3.312  | 0.000925 *** |
| V4A410      | -1.667e+00 | 9.565e-01  | -1.743  | 0.081297 .   |
| V4A42       | -6.392e-01 | 3.202e-01  | -1.996  | 0.045894 *   |
| V4A43       | -7.132e-01 | 3.007e-01  | -2.372  | 0.017680 *   |
| V4A44       | 4.628e-01  | 1.009e+00  | 0.459   | 0.646586     |
| V4A45       | 6.217e-01  | 6.758e-01  | 0.920   | 0.357584     |
| V4A46       | 1.646e-01  | 4.795e-01  | 0.343   | 0.731397     |

|         |            |           |        |             |
|---------|------------|-----------|--------|-------------|
| V4A48   | -1.501e+01 | 5.061e+02 | -0.030 | 0.976335    |
| V4A49   | -8.667e-01 | 4.136e-01 | -2.096 | 0.036115 *  |
| V5      | 6.817e-05  | 5.545e-05 | 1.230  | 0.218876    |
| V6A62   | -1.963e-01 | 3.660e-01 | -0.536 | 0.591674    |
| V6A63   | -1.036e+00 | 5.281e-01 | -1.961 | 0.049875 *  |
| V6A64   | -1.468e+00 | 6.067e-01 | -2.420 | 0.015540 *  |
| V6A65   | -8.774e-01 | 3.113e-01 | -2.818 | 0.004827 ** |
| V7A72   | 1.205e-01  | 5.240e-01 | 0.230  | 0.818069    |
| V7A73   | -2.039e-01 | 4.980e-01 | -0.409 | 0.682242    |
| V7A74   | -1.006e+00 | 5.471e-01 | -1.838 | 0.066061 .  |
| V7A75   | -1.709e-01 | 5.073e-01 | -0.337 | 0.736167    |
| V8      | 2.737e-01  | 1.081e-01 | 2.532  | 0.011338 *  |
| V9A92   | -8.649e-01 | 4.723e-01 | -1.831 | 0.067040 .  |
| V9A93   | -1.136e+00 | 4.626e-01 | -2.457 | 0.014017 *  |
| V9A94   | -9.721e-01 | 5.649e-01 | -1.721 | 0.085277 .  |
| V10A102 | 3.078e-01  | 5.033e-01 | 0.612  | 0.540792    |
| V10A103 | -6.958e-01 | 4.863e-01 | -1.431 | 0.152462    |
| V11     | -8.948e-03 | 1.055e-01 | -0.085 | 0.932411    |
| V12A122 | 5.914e-01  | 3.098e-01 | 1.909  | 0.056232 .  |
| V12A123 | 3.720e-01  | 2.982e-01 | 1.247  | 0.212255    |
| V12A124 | 9.804e-01  | 5.343e-01 | 1.835  | 0.066518 .  |
| V13     | -1.435e-02 | 1.109e-02 | -1.294 | 0.195741    |
| V14A142 | -1.444e-01 | 4.823e-01 | -0.299 | 0.764668    |
| V14A143 | -6.105e-01 | 2.995e-01 | -2.038 | 0.041524 *  |
| V15A152 | -6.422e-01 | 2.876e-01 | -2.233 | 0.025529 *  |
| V15A153 | -1.035e+00 | 5.906e-01 | -1.753 | 0.079533 .  |
| V16     | 1.728e-01  | 2.369e-01 | 0.729  | 0.465829    |
| V17A172 | 2.852e-01  | 1.004e+00 | 0.284  | 0.776407    |
| V17A173 | 5.242e-01  | 9.678e-01 | 0.542  | 0.588097    |

```
V17A174  3.168e-01 9.668e-01 0.328 0.743133
V18      -1.682e-02 3.082e-01 -0.055 0.956461
V19A192  -1.709e-01 2.427e-01 -0.704 0.481255
V20A202  -1.418e+00 7.042e-01 -2.014 0.044048 *
---
Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

(Dispersion parameter for binomial family taken to be 1)

```
Null deviance: 841.21 on 699 degrees of freedom
Residual deviance: 611.85 on 651 degrees of freedom
AIC: 709.85
```

Number of Fisher Scoring iterations: 14

#Based on above Test, we identify insignificant variables based on the level or variables p-value >0.05.  
#Drop V20, V19, V18, V17, V15, V13, V12, V11, V10, V7, V5 since those terms had p-values greater than 0.05.

#Check for multicollinearity

#

```
library(car)
```

```
vif(reg1)
```

```
alias(reg1)
```

#

#Variables V2 and V5 have slightly elevated levels of variable inflation factors.

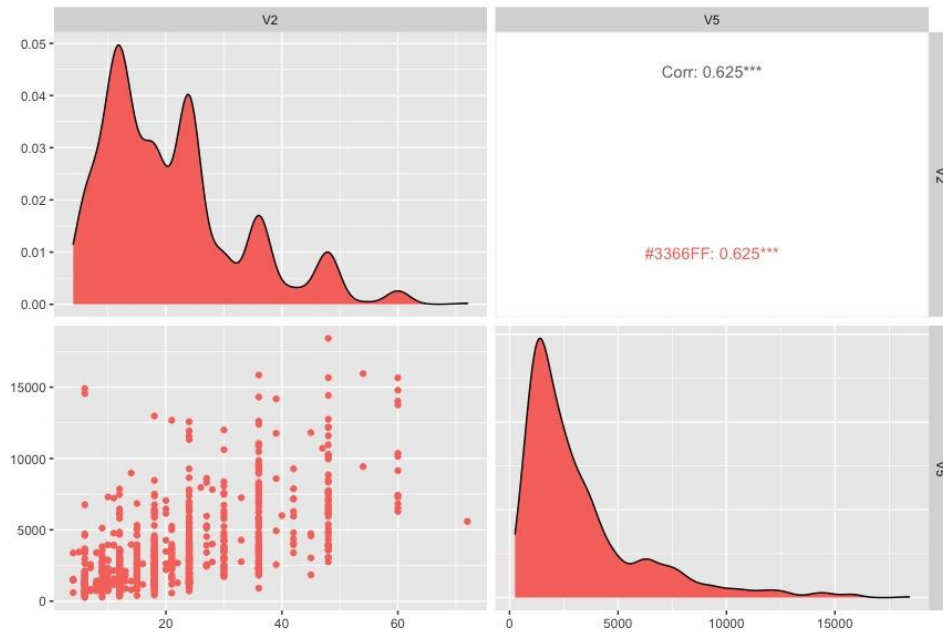
#There could be a positive relationship between duration of the loan and credit amount since bigger loans have larger loan durations.

#Use scatter plot to explore the extent

#

```
library(GGally)
```

```
ggpairs(data_gc, columns = c('V2', 'V5'), mapping=ggplot2::aes(color= '#3366FF'))
```



#

#confirms correlation is positive. The anova test revealed we could drop V5 so we'll just retain V2.

#Rerun the model without the insignificant predictors V20, V19, V18, V17, V15, V13, V12, V11, V10, V7, V5 from the model and rerun

#

```
reg2<-update(reg1,~.-V20-V19-V18-V17-V15-V13-V12-V11-V10-V7-V5)
```

```
summary(reg2)
```

**Call:**

```
glm(formula = V21 ~ V1 + V2 + V3 + V4 + V6 + V8 + V9 + V14 +  
      V16, family = binomial(link = "logit"), data = builddata)
```

**Deviance Residuals:**

| Min     | 1Q      | Median  | 3Q     | Max    |
|---------|---------|---------|--------|--------|
| -2.0997 | -0.7219 | -0.4122 | 0.7047 | 2.7850 |

Coefficients:

|             | Estimate   | Std. Error | z value | Pr(> z )     |
|-------------|------------|------------|---------|--------------|
| (Intercept) | 0.899443   | 0.833958   | 1.079   | 0.28080      |
| V1A12       | -0.229943  | 0.248998   | -0.923  | 0.35576      |
| V1A13       | -1.479960  | 0.506436   | -2.922  | 0.00347 **   |
| V1A14       | -1.538294  | 0.265202   | -5.800  | 6.61e-09 *** |
| V2          | 0.044739   | 0.008459   | 5.289   | 1.23e-07 *** |
| V3A31       | -0.367203  | 0.623662   | -0.589  | 0.55601      |
| V3A32       | -0.881794  | 0.486126   | -1.814  | 0.06969 .    |
| V3A33       | -0.754188  | 0.539185   | -1.399  | 0.16189      |
| V3A34       | -1.568869  | 0.492545   | -3.185  | 0.00145 **   |
| V4A41       | -1.278108  | 0.428456   | -2.983  | 0.00285 **   |
| V4A410      | -1.248226  | 0.885577   | -1.410  | 0.15869      |
| V4A42       | -0.311196  | 0.296121   | -1.051  | 0.29330      |
| V4A43       | -0.740112  | 0.280919   | -2.635  | 0.00842 **   |
| V4A44       | 0.739196   | 0.952095   | 0.776   | 0.43752      |
| V4A45       | 0.568370   | 0.643672   | 0.883   | 0.37723      |
| V4A46       | 0.362919   | 0.451306   | 0.804   | 0.42131      |
| V4A48       | -15.078590 | 505.775652 | -0.030  | 0.97622      |
| V4A49       | -0.799780  | 0.389145   | -2.055  | 0.03986 *    |
| V6A62       | -0.074669  | 0.344678   | -0.217  | 0.82849      |
| V6A63       | -1.006593  | 0.511247   | -1.969  | 0.04897 *    |
| V6A64       | -1.439070  | 0.579647   | -2.483  | 0.01304 *    |
| V6A65       | -0.897226  | 0.295765   | -3.034  | 0.00242 **   |
| V8          | 0.201132   | 0.091310   | 2.203   | 0.02761 *    |
| V9A92       | -0.632101  | 0.438133   | -1.443  | 0.14910      |
| V9A93       | -1.099886  | 0.428547   | -2.567  | 0.01027 *    |
| V9A94       | -0.877724  | 0.537281   | -1.634  | 0.10233      |



V14A142 -0.084506 0.458866 -0.184 0.85389

V14A143 -0.547113 0.284601 -1.922 0.05456 .

V16 0.151009 0.219033 0.689 0.49055

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 841.21 on 699 degrees of freedom

Residual deviance: 646.94 on 671 degrees of freedom

AIC: 704.94

Number of Fisher Scoring iterations: 14

#

#We now have a new model with selected terms. AIC has improved for the model since it decreased from 709 to 704.94

#The 4th variable has a really strange coefficient estimate for A48 and very high p-values for both A44 and A48.

#Need to explore this more and decide whether to combine these levels with another level.

#

View(builddata\$V4)

#

#After viewing variable 4 which is a categorical variable for the Purpose of the loan I found there is only 4 observations for A48 factor level which stands for the loan purpose of retraining.

#This purpose has such a low number of observations in the build set and is causing instability in its estimation so I'll combine A48 with the base level A40 (car new) and rerun the model.

#Factor A44 has twice as many observation in the build set but at 8 that number is still low, so I will also combine this level with the base to zero out the coefficient.

#See if AIC improves.

#

```

levels(builddata$V4)[levels(builddata$V4)==c("A48","A44")]<-"A40"
reg2.5<-update(reg2,~.)
summary(reg2.5)

#
#AIC value is not changed much... so improved the stability of the model by relabelling the variable.
#Show the cross validated performance of the model Using 7 folds so we have 100 obs in each fold
#
library(caret)
builddata$risk<-as.factor(ifelse(builddata$V21==1,'bad','good'))
buildfolds<-caret::createFolds(builddata$risk,k=7)
set.seed(478)
cvrm<-train(risk ~ V1 + V2 + V3 + V4 + V6 + V8 + V9 + V14 + V16,
            data=builddata,
            method='glm',
            trControl=trainControl(method = 'cv', number = 7, index = buildfolds, classProbs = TRUE,
            summaryFunction = twoClassSummary),
            metric='ROC')
cvrm

```

Generalized Linear Model

**700 samples**

**9 predictor**

**2 classes: 'bad', 'good'**

**No pre-processing**

**Resampling: Cross-Validated (7 fold)**

**Summary of sample sizes: 100, 100, 100, 100, 100, 100, ...**

### Resampling results:

| ROC       | Sens      | Spec      |
|-----------|-----------|-----------|
| 0.6530515 | 0.4018385 | 0.7988784 |

```
cvrn$finalModel
```

```
#
```

```
#Here is the 7-fold cross validated area under the curve=0.6788661, sensitivity=0.4526674 and  
specificity= 0.8099621
```

```
#assuming a 50% threshold and '1' is the positive class.
```

```
#Caret found some predictions were equal to 0 or 1 when the output is probability which happens with  
logit regression.
```

```
#performance changes on the test set for the 50% threshold.
```

```
#Make prediction on the test set similar to Buildata
```

```
#
```

```
levels(testdata$V4)[levels(testdata$V4)==c("A48","A44")]<-"A40"
```

```
reg2test<-predict(reg2.5, newdata=testdata,type='response')
```

```
#set threshold at 50%
```

```
reg2testfact<-as.factor(ifelse(reg2test>0.5,1,0))#positive class is bad risks
```

```
confusionMatrix(reg2testfact,as.factor(testdata$V21))#using 50% threshold
```

### Confusion Matrix and Statistics

#### Reference

| Prediction | 0 | 1 |
|------------|---|---|
|------------|---|---|

|   |     |    |
|---|-----|----|
| 0 | 178 | 51 |
|---|-----|----|

|   |    |    |
|---|----|----|
| 1 | 24 | 47 |
|---|----|----|

**Accuracy : 0.75**

**95% CI : (0.697, 0.798)**

**No Information Rate : 0.6733**

**P-Value [Acc > NIR] : 0.002357**

**Kappa : 0.3883**

**Mcnemar's Test P-Value : 0.002680**

**Sensitivity : 0.8812**

**Specificity : 0.4796**

**Pos Pred Value : 0.7773**

**Neg Pred Value : 0.6620**

**Prevalence : 0.6733**

**Detection Rate : 0.5933**

**Detection Prevalence : 0.7633**

**Balanced Accuracy : 0.6804**

**'Positive' Class : 0**

```
library(pROC)
```

```
roc(testdata$V21,ifelse(reg2test>0.5,1,0))#inputs must be numeric
```

```
plot.roc(testdata$V21,ifelse(reg2test>0.5,1,0))
```

```
# Total Accuracy 75%
```

```
#Sensitivity 88%
```

```
# Specificity 47%
```

---

**10.3.2. Because the model gives a result between 0 and 1, it requires setting a threshold probability to separate between “good” and “bad” answers. In this data set, they estimate that incorrectly identifying a bad customer as good, is 5 times worse than incorrectly classifying a good customer as bad. Determine a good threshold probability based on your model.**

```
costs = matrix(c(0, 5, 1, 0), nrow = 2)
dimnames(costs) = list(Actual = c("good", "bad"), Predicted= c("good", "bad"))
print(costs) #Cost of misclassifications 5x for FP vs. FN
```

```
      Predicted
Actual good bad
good    0  1
bad     5  0
```

```
#initialize list
```

```
cost <- vector(mode = "list")
```

```
set.seed(713)
```

```
predicted <- predict(reg2 , testdata, type="response")
```

```
for (i in 1:100){
```

```
  predicted_roundup <- as.integer(predicted > i/100 )
```

```
  cm_matrix <- as.matrix(table(testdata$V21 ,predicted_roundup))
```

```
  #Ensuring NO out of bounds issues while looping
```

```

if(nrow(cm_matrix)==2) {fp<-cm_matrix[2,1]} else {fp=0}
if(ncol(cm_matrix)==2){fn<-cm_matrix[1,2]} else {fn=0}

cost<-c(cost, fn*1+fp*5)
}

#Plots ov Total cost vs % thresholds

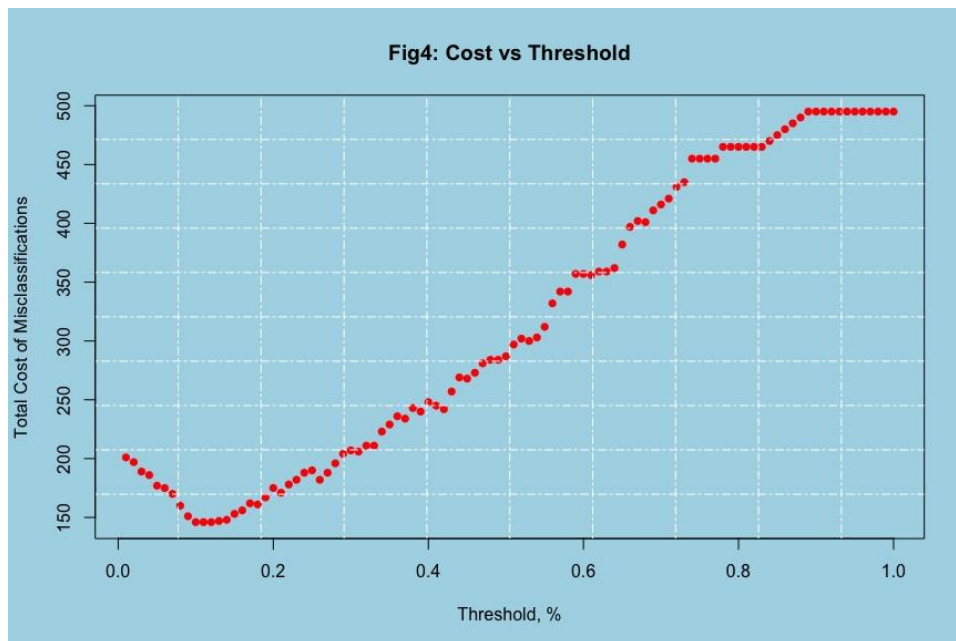
par(bg = 'lightblue')

plot(x=seq(0.01,1,by=0.01),y=cost,xlab = "Threshold, %",ylab = "Total Cost of Misclassifications",main =
"Fig4: Cost vs Threshold",col='red',pch=16)

grid (10,10, lty = 6, col = "white")

numerator<-which.min(cost)
min.threshold <-numerator/100
print(sprintf("Minimum Threshold due to Uneven Error Cost= %0.3f", min.threshold))

```



```

# "Minimum Threshold due to Uneven Error Cost= 0.100"

```

